

1. Import Necessary Libraries

```
In [1]: import pandas as pd
```

2. Import Dataset

```
In [2]: claimants_data=pd.read_csv('claimants.csv')  
claimants_data
```

```
Out[2]:
```

	CASENUM	ATTORNEY	CLMSEX	CLMINSUR	SEATBELT	CLMAGE	LOSS
0	5	0	0.0	1.0	0.0	50.0	34.940
1	3	1	1.0	0.0	0.0	18.0	0.891
2	66	1	0.0	1.0	0.0	5.0	0.330
3	70	0	0.0	1.0	1.0	31.0	0.037
4	96	1	0.0	1.0	0.0	30.0	0.038
...
1335	34100	1	0.0	1.0	0.0	NaN	0.576
1336	34110	0	1.0	1.0	0.0	46.0	3.705
1337	34113	1	1.0	1.0	0.0	39.0	0.099
1338	34145	0	1.0	0.0	0.0	8.0	3.177
1339	34153	1	1.0	1.0	0.0	30.0	0.688

1340 rows × 7 columns

```
In [3]: claimants_data['ATTORNEY'].unique()
```

```
Out[3]: array([0, 1])
```

3. Data Understanding

3.1 Initial Analysis

```
In [4]: claimants_data.shape
```

```
Out[4]: (1340, 7)
```

```
In [5]: claimants_data.isna().sum()
```

```
Out[5]: CASENUM      0  
ATTORNEY      0  
CLMSEX       12  
CLMINSUR     41  
SEATBELT      48  
CLMAGE      189  
LOSS          0  
dtype: int64
```

```
In [6]: claimants_data.dtypes
```

```
Out[6]: CASENUM      int64  
ATTORNEY      int64  
CLMSEX       float64  
CLMINSUR     float64  
SEATBELT      float64  
CLMAGE       float64  
LOSS         float64  
dtype: object
```

4. Data Preparation

```
In [7]: del claimants_data['CASENUM']
```

```
In [8]: claimants_data.head()
```

```
Out[8]: ATTORNEY CLMSEX CLMINSUR SEATBELT CLMAGE LOSS
0 0 0.0 1.0 0.0 50.0 34.940
1 1 1.0 0.0 0.0 18.0 0.891
2 1 0.0 1.0 0.0 5.0 0.330
3 0 0.0 1.0 1.0 31.0 0.037
4 1 0.0 1.0 0.0 30.0 0.038
```

```
In [9]: claimants_data.dropna(axis=0,inplace=True)
```

```
In [10]: claimants_data.isna().sum()
```

```
Out[10]: ATTORNEY 0
CLMSEX 0
CLMINSUR 0
SEATBELT 0
CLMAGE 0
LOSS 0
dtype: int64
```

```
In [11]: claimants_data.shape
```

```
Out[11]: (1096, 6)
```

```
In [12]: #Seeing balanced or unbalanced class
claimants_data['ATTORNEY'].value_counts()
```

```
Out[12]: ATTORNEY
0 578
1 518
Name: count, dtype: int64
```

5.model Building

2 steps process:

1.Separate Input and Output 2.Perform some Model Validation Techniques:

- *Train-Test Split
- *K-Fold Cross Validation
- *Leave One Out Cross Validation

```
In [13]: X = claimants_data.drop('ATTORNEY',axis=1)
Y = claimants_data[['ATTORNEY']]
```

```
In [14]: X.shape,Y.shape
```

```
Out[14]: ((1096, 5), (1096, 1))
```

```
In [15]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(X,Y,test_size=0.20,random_state=123,shuffle=True)
```

```
In [16]: x_train.shape,y_train.shape
```

```
Out[16]: ((876, 5), (876, 1))
```

```
In [17]: x_test.shape,y_test.shape
```

```
Out[17]: ((220, 5), (220, 1))
```

6.Model Training

```
In [18]: import warnings
warnings.filterwarnings('ignore')
```

```
In [19]: from sklearn.linear_model import LogisticRegression
logistic_model = LogisticRegression()
logistic_model.fit(x_train,y_train)
```

```
Out[19]: ▾ LogisticRegression ⓘ ⓘ  
LogisticRegression()
```

```
In [20]: # from sklearn.tree import DecisionTreeClassifier  
# dt = DecisionTreeClassifier()  
# dt.fit(x_train,y_train)
```

```
In [21]: %%time  
logistic_model.fit(x_train,y_train)  
CPU times: total: 31.2 ms  
Wall time: 57.1 ms  
Out[21]: ▾ LogisticRegression ⓘ ⓘ  
LogisticRegression()
```

7. Model Testing

```
In [22]: from sklearn.metrics import accuracy_score,confusion_matrix,classification_report,roc_auc_score,roc_curve
```

```
In [23]: y_pred_train=logistic_model.predict(x_train)
```

```
In [24]: accuracy_score(y_train,y_pred_train)
```

```
Out[24]: 0.7134703196347032
```

```
In [25]: y_pred_test=logistic_model.predict(x_test)  
accuracy_score(y_test,y_pred_test)
```

```
Out[25]: 0.6863636363636364
```

```
In [26]: confusion_matrix(y_train,y_pred_train)
```

```
Out[26]: array([[315, 151],  
                 [100, 310]])
```

```
In [27]: #Accuracy  
(315+310)/(315+151+100+310)
```

```
Out[27]: 0.7134703196347032
```

```
In [28]: #Recall Zero row wise  
rzero = 315/(315+151)  
rzero
```

```
Out[28]: 0.6759656652360515
```

```
In [29]: #Recall One  
rone = 310/(100+310)  
rone
```

```
Out[29]: 0.7560975609756098
```

```
In [30]: #Precision zero  
pzero = 315/(315+100)  
pzero
```

```
Out[30]: 0.7590361445783133
```

```
In [31]: #Precision one  
pone = 310/(151+310)  
pone
```

```
Out[31]: 0.6724511930585684
```

```
In [32]: avg = (0.6759656652360515+0.7560975609756098)/2  
avg
```

```
Out[32]: 0.7160316131058306
```

```
In [33]: avgp = (0.7590361445783133+0.6724511930585684)/2  
avgp
```

```
Out[33]: 0.7157436688184409
```

```
In [34]: #F1 Score  
f1 = (2*avgr*avgp)/(avgr+avgp)  
f1
```

```
Out[34]: 0.715887612007902
```


```
In [35]: #Direct function  
print(classification_report(y_train,y_pred_train))
```

	precision	recall	f1-score	support
0	0.76	0.68	0.72	466
1	0.67	0.76	0.71	410
accuracy			0.71	876
macro avg	0.72	0.72	0.71	876
weighted avg	0.72	0.71	0.71	876

```
In [36]: from sklearn.metrics import roc_curve,roc_auc_score  
fpr, tpr, thresholds = roc_curve(Y,logistic_model.predict_proba(X)[:,1])  
  
auc = roc_auc_score(y_train,y_pred_train)  
print(auc)  
  
import matplotlib.pyplot as plt  
plt.plot(fpr, tpr, color = 'red', label = 'logit model(area =%0.2f)'%auc)  
plt.plot([0,1],[0,1],'k--')  
plt.xlabel('False Positive Rate or [1 - True Negative Rate]')  
plt.ylabel('True Positive Rate')
```

```
0.7160316131058306
```

```
Out[36]: Text(0, 0.5, 'True Positive Rate')
```



```
In [37]: from sklearn.neighbors import KNeighborsClassifier
```

```
In [38]: knn = KNeighborsClassifier() #build model  
knn.fit(x_train,y_train)  
y_pred_train= knn.predict(x_train)  
y_pred_test= knn.predict(x_test)
```

```
In [39]: confusion_matrix(y_train,y_pred_train)
```

```
Out[39]: array([[361, 105],  
 [ 98, 312]])
```

```
In [40]: print(classification_report(y_train,y_pred_train))
```

```

precision    recall  f1-score   support

          0       0.79      0.77      0.78      466
          1       0.75      0.76      0.75      410

   accuracy                           0.77      876
macro avg       0.77      0.77      0.77      876
weighted avg    0.77      0.77      0.77      876

```

In [41]: `confusion_matrix(y_test,y_pred_test)`

Out[41]: `array([[67, 45],
 [38, 70]])`

In [42]: `print(classification_report(y_test,y_pred_test))`

```

precision    recall  f1-score   support

          0       0.64      0.60      0.62      112
          1       0.61      0.65      0.63      108

   accuracy                           0.62      220
macro avg       0.62      0.62      0.62      220
weighted avg    0.62      0.62      0.62      220

```

In [43]: `# how to find best optimal k value`

In []:

8. Test Accuracy

In [44]: `y_pred_test=logistic_model.predict(x_test)`

In [45]: `accuracy_score(y_test,y_pred_test)
print(classification_report(y_test,y_pred_test))`

```

precision    recall  f1-score   support

          0       0.74      0.59      0.66      112
          1       0.65      0.79      0.71      108

   accuracy                           0.69      220
macro avg       0.70      0.69      0.68      220
weighted avg    0.70      0.69      0.68      220

```

In [46]: `from sklearn.ensemble import AdaBoostClassifier
from sklearn.tree import DecisionTreeClassifier`

In [47]: `base_model = DecisionTreeClassifier(max_depth=1)
abm = AdaBoostClassifier(estimator=base_model,n_estimators=1000,random_state=32,learning_rate=1.0)
abm.fit(x_train,y_train)`

Out[47]:

```

▶ AdaBoostClassifier
  ▶ estimator:
    DecisionTreeClassifier
      ▶ DecisionTreeClassifier

```

In [48]: `y_pred_train = abm.predict(x_train)
print(classification_report(y_train,y_pred_train))`

```

precision    recall  f1-score   support

          0       0.75      0.78      0.76      466
          1       0.74      0.70      0.72      410

   accuracy                           0.74      876
macro avg       0.74      0.74      0.74      876
weighted avg    0.74      0.74      0.74      876

```

In [49]: `y_pred_test = abm.predict(x_test)
print(classification_report(y_test,y_pred_test))`

```

precision    recall   f1-score   support
0            0.69     0.80      0.74      112
1            0.76     0.63      0.69      108

accuracy                           0.72      220
macro avg       0.72     0.72      0.72      220
weighted avg    0.72     0.72      0.72      220

```

```
In [50]: from sklearn.ensemble import GradientBoostingClassifier
```

```

In [51]: base_model = DecisionTreeClassifier()
model = GradientBoostingClassifier()
model.fit(x_train,y_train)
y_pred_train = model.predict(x_train)
y_pred_test = model.predict(x_test)

```

```
In [52]: print(classification_report(y_train,y_pred_train))
```

	precision	recall	f1-score	support
0	0.79	0.83	0.81	466
1	0.80	0.76	0.78	410
accuracy			0.80	876
macro avg	0.80	0.79	0.79	876
weighted avg	0.80	0.80	0.80	876

```
In [53]: print(classification_report(y_test,y_pred_test))
```

	precision	recall	f1-score	support
0	0.67	0.76	0.71	112
1	0.71	0.62	0.66	108
accuracy			0.69	220
macro avg	0.69	0.69	0.69	220
weighted avg	0.69	0.69	0.69	220

9. Model Deployment

```
In [56]: # from pickle import dump
```

```
In [57]: # dump(logistic_model,open('claimants_intelligene.pkl','wb'))
```

```
In [58]: # from pickle import load
```

```
In [59]: # loaded_intelligence=load(open('claimants_intelligene.pkl','rb'))
```

```
In [60]: # y_pred_deployment=loaded_intelligence.predict(x_test)
```

1. STANDARDIZE AND CHECK THE ACCURACY

```

In [61]: from sklearn.preprocessing import StandardScaler
std_scaler=StandardScaler()
scaled_x=std_scaler.fit_transform(X)
scaled_x

```

```

Out[61]: array([[-1.13916369,  0.32550512, -0.13633547,  1.05048704,  2.96924493],
   [ 0.87783697, -3.07214831, -0.13633547, -0.51942439, -0.28328699],
   [-1.13916369,  0.32550512, -0.13633547, -1.15720091, -0.33687653],
   ...,
   [ 0.87783697,  0.32550512, -0.13633547,  0.51082998, -0.35894281],
   [ 0.87783697, -3.07214831, -0.13633547, -1.01002171, -0.06491676],
   [ 0.87783697,  0.32550512, -0.13633547,  0.0692924 , -0.30267857]])

```

```
In [62]: x_train,x_test,y_train,y_test=train_test_split(scaled_x,Y,test_size=0.20,random_state=123)
```

```

In [63]: %time
logistic_model.fit(x_train,y_train)

```

CPU times: total: 15.6 ms
Wall time: 64 ms

```
Out[63]:
```

```
▼ LogisticRegression ⓘ ⓘ  
LogisticRegression()
```

```
In [64]: y_pred_train = logistic_model.predict(x_train)
```

```
In [65]: accuracy_score(y_train,y_pred_train)
```

```
Out[65]: 0.7180365296803652
```

```
In [66]: y_pred_test = logistic_model.predict(x_test)
```

```
In [67]: accuracy_score(y_test,y_pred_test)
```

```
Out[67]: 0.66818181818182
```