

1) Import Necessary Libraries

```
In [49]: import pandas as pd
from matplotlib import pyplot as plt
import seaborn as sns

import warnings
warnings.filterwarnings('ignore')
```

2) Import Data

```
In [50]: cars_data=pd.read_csv('Cars.csv')
cars_data
```

```
Out[50]:   HP      MPG      VOL       SP       WT
  0  49  53.700681    89  104.185353  28.762059
  1  55  50.013401   92  105.461264  30.466833
  2  55  50.013401   92  105.461264  30.193597
  3  70  45.696322   92  113.461264  30.632114
  4  53  50.504232   92  104.461264  29.889149
 ...
  76 322  36.900000   50  169.598513  16.132947
  77 238  19.197888  115  150.576579  37.923113
  78 263  34.000000   50  151.598513  15.769625
  79 295  19.833733  119  167.944460  39.423099
  80 236  12.101263  107  139.840817  34.948615
```

81 rows × 5 columns

3) Data Understanding

3.1 Initial Analysis

```
In [4]: cars_data.shape
```

```
Out[4]: (81, 5)
```

```
In [5]: cars_data.isna().sum()
```

```
Out[5]: HP      0
MPG      0
VOL      0
SP       0
WT       0
dtype: int64
```

```
In [6]: cars_data.dtypes
```


```
Out[6]: HP      int64
MPG     float64
VOL      int64
SP      float64
WT      float64
dtype: object
```

3.2 Assumptions Check

Assumption 1: Test For Linearity


```
In [7]: sns.scatterplot(x='HP',y='MPG',data=cars_data)
plt.title('HP Vs MPG')
plt.show()
```

HP Vs MPG




```
In [8]: sns.lmplot(x='VOL',y='MPG',data=cars_data)
plt.title('VOL Vs MPG')
plt.show()
```

VOL Vs MPG




```
In [10]: sns.lmplot(x='SP',y='MPG',data=cars_data)
plt.title('SP Vs MPG')
plt.show()
```

SP Vs MPG




```
In [12]: sns.lmplot(x='WT', y='MPG', data=cars_data)
plt.title('WT VS MPG')
plt.show()
```

WT VS MPG



```
In [14]: sns.pairplot(data=cars_data)
plt.show()
```



In [15]: `cars_data.corr()`

Out[15]:

	HP	MPG	VOL	SP	WT
HP	1.000000	-0.725038	0.077459	0.973848	0.076513
MPG	-0.725038	1.000000	-0.529057	-0.687125	-0.526759
VOL	0.077459	-0.529057	1.000000	0.102170	0.999203
SP	0.973848	-0.687125	0.102170	1.000000	0.102439
WT	0.076513	-0.526759	0.999203	0.102439	1.000000


Observation

The linearity test failed


Asssumptions 2: Test for Linearity

2.1 : Using Distplot


In [9]: `sns.distplot(a=cars_data['HP'], hist=False)
plt.show()`




```
In [10]: sns.distplot(a=cars_data['VOL'],hist=False)
plt.show()
```



```
In [11]: sns.distplot(a=cars_data['SP'],hist=False)
plt.show()
```



```
In [12]: sns.distplot(a=cars_data['WT'],hist=False)
plt.show()
```




Observation


Normality test failed

2.2:Using Probplot


```
In [15]: from scipy import stats
stats.probplot(x=cars_data['HP'],dist='norm',plot=plt)
plt.show()
```




```
In [16]: stats.probplot(x=cars_data['VOL'],dist='norm',plot=plt)  
plt.show()
```



```
In [17]: stats.probplot(x=cars_data['SP'],dist='norm',plot=plt)  
plt.show()
```



```
In [18]: stats.probplot(x=cars_data['WT'], dist='norm', plot=plt)
plt.show()
```



Observation

Test for linearity failed

Assumption 3: Test for Multi collinearity

By using 2 Techniques

1. Correlation Matrix

2. Variance Inflation Factor (VIF)

3.1 Correlation Test

```
In [ ]: corr_matrix=cars_data.corr().round()
corr_matrix
```

```
In [ ]: sns.heatmap(data=corr_matrix, annot=True)
plt.show()
```

Observation

There is multicollinearity in my data, so this test also fails.

4 No Auto Regression Passed

Model Building

```
In [25]: #x=cars_data.drop(labels='MPG',axis=1,inplace=True)
x=cars_data[['HP','VOL','SP','WT']]
y=cars_data['MPG']
```

cars_data.drop(['MPG'],axis=1)

```
In [26]: x
```

```
Out[26]:   HP  VOL      SP      WT
0    49    89  104.185353  28.762059
1    55    92  105.461264  30.466833
2    55    92  105.461264  30.193597
3    70    92  113.461264  30.632114
4    53    92  104.461264  29.889149
...
76   322   50  169.598513  16.132947
77   238   115  150.576579  37.923113
78   263   50  151.598513  15.769625
79   295   119  167.944460  39.423099
80   236   107  139.840817  34.948615
```

81 rows × 4 columns

```
In [27]: y
```

```
Out[27]: 0      53.700681
1      50.013401
2      50.013401
3      45.696322
4      50.504232
...
76     36.900000
77     19.197888
78     34.000000
79     19.833733
80     12.101263
Name: MPG, Length: 81, dtype: float64
```

Model Validation / Data Validation

```
In [33]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,random_state=35)
```

```
In [34]: print(x_train.shape,x_test.shape,y_train.shape,y_test.shape)
(64, 4) (17, 4) (64,) (17,)
```

```
In [35]: from sklearn.linear_model import LinearRegression
```

```
In [37]: mlr = LinearRegression() #model Building
mlr.fit(x_train,y_train)
y_pred_train=mlr.predict(x_train)
y_pred_test=mlr.predict(x_test)
```

```
In [ ]: x1_train= hp - mpg
x2_train=hp,vol - mpg
x3_train=hp,vol,sp - mpg
```

```
In [39]: y_pred_test
```

```
Out[39]: array([41.93599129, 36.64877574, 37.21897411, 33.80175655, 39.2043074 ,  
   34.13391752, 35.61010219, 38.83215133, 35.77879072, 33.53455187,  
   23.02614465, 43.54665895, 11.49873125, 34.36553633, 21.54290653,  
   24.03684324, 48.98206504])
```

```
In [40]: error_train=y_train-y_pred_train  
error_train
```

```
Out[40]: 14      3.440987  
8       -2.879342  
47      -0.711218  
52      0.889579  
49      -4.197923  
...  
63      -3.055257  
33      1.492869  
55      -0.136305  
15      2.445640  
73      0.577655  
Name: MPG, Length: 64, dtype: float64
```

```
In [45]: error_test=y_test-y_pred_test  
error_test
```

```
Out[45]: 13      2.716843  
48      -5.634645  
22      1.091632  
51      -4.171821  
38      -5.133639  
46      0.427581  
41      -0.457375  
27      -0.421148  
58      -5.646868  
50      -3.904616  
67      0.077027  
0       10.154022  
79      8.335002  
40      0.787191  
71      1.660662  
68      -0.933672  
7       -2.265511  
Name: MPG, dtype: float64
```

```
In [51]: from sklearn.metrics import r2_score  
r2_train=r2_score(y_train,y_pred_train)  
print(r2_train)
```

```
0.7689342881829975
```

```
In [52]: r2_test=r2_score(y_test,y_pred_test)  
print(r2_test)
```

```
0.7618997645852453
```